


to revise:

- notes
- labs
- revision guide
- Piazza

FNLP

final exam notes



} syntax

→ part of speech

this is a simple sentence
 be 3sg present SIMPLE1 having few parts SENTENCE1 string of words satisfying the grammatical rules of a language

→ words

→ morphology

→ semantics

why is NLP hard?


• NLP is hard b/c of ambiguity But it is an instructive one → discourse

- word senses: bank (finance or river?)
- part of speech: chair (noun or verb?)
- syntactic structure: I saw a man with a telescope
- quantifier scope: Every child loves some movie
- multiple: I saw her duck.

• data is sparse due to Zipf's law: freq vs. rank

$f \cdot r \approx k$ where f = freq of word, r = rank (by f), k = constant

$$\Rightarrow \log f = \log k - \log r \quad y = b + mx$$



• variation: think WSJ vs. social media

• expressivity: can express same thing many different ways

• context dependence and unknown representation: need world knowledge (sometimes/usually), don't know how to represent "meaning"

(2) text corpora

- good if naturally-occurring, includes metadata, or has linguistic annotations
- for sentiment analysis, good to have some sort of ratings/labels
- problems w/ counting positive/negative words: sense ambiguity, sarcasm/irony, context (other)
- tokenization: adding logical boundaries between separate word/punctuation tokens not already separated by spaces
 - only one part of preprocessing/normalization
- "domain" of corpus includes mode of communication (e.g., speech, writing), topic (e.g., politics, physics), genre (e.g., news, tweet), audience (formality, politeness, complexity)
- domain adaption corrects for when training data and test data don't come from the same source

(3) n-gram language models

- this lecture talks about sentence probabilities
- language model gives approx. for true prob. of an arbitrary sequence of words
- uses of a language model: spelling correction, automatic speech recognition, machine translation
 - in each case, get possibilities from some model (e.g., error model, acoustic model, translation model) and use language model to choose among alternatives
- unrealistic to do MLE on complete sentences; instead, multiply probabilities of sub-parts
- n-gram model:

$$P(w_1, \dots, w_n) = \prod_{i=1}^n P(w_i | w_1, w_2, \dots, w_{i-1}) \approx P(w_i | w_{i-2}, w_{i-1}) \text{ for trigram model} \\ \approx P(w_i | w_{i-1}) \text{ for bigram} \\ \approx P(w_i) \text{ for unigram}$$

• estimation w/ trigram model:

$$P_{MLE}(w_i | w_{i-2}, w_{i-1}) = \frac{C(w_{i-2}, w_{i-1}, w_i)}{C(w_{i-2}, w_{i-1})}$$

n-gram language models (cont'd)

- practical implementation of tri-grams:

$$P(\vec{w}) = P(w_1) P(w_2 | w_1) \prod_{i=3}^n P(w_i | w_{i-2}, w_{i-1})$$

- but this doesn't capture the fact that some word strings are more/less likely to start/end a sentence
- can assume $w_{-1} = w_0 = \langle s \rangle$ and $w_{n+1} = \langle s \rangle$;

$$P(\vec{w}) = \prod_{i=1}^{n+1} P(w_i | w_{i-2}, w_{i-1})$$

- typically use negative log probs since probs can get small

Language Models: Smoothing and evaluation

- ways of choosing between models (e.g., bigram and trigram)

- extrinsic: plug into downstream system and (somehow) evaluate that system
- intrinsic: inherent to the current task; usually quicker/easier

- one intrinsic measure: entropy:

$$H(X) = \sum_x -P(x) \log_2 P(x) = \mathbb{E}[-\log_2 P(X)]$$

- this measures uncertainty / disorder / surprise



$$H(X)=0$$



$$H(X)=1$$



$$H(X)=2$$



$$H(X)=1.4$$



$$H(X)=0.2$$

or, # yes/no questions

- intuition: average number of bits needed to encode $X \hat{=}$ entropy of X

- a good model assigns high probability to sequences that actually have high probability; put another way, our model should have low uncertainty (entropy) about which word comes next \Rightarrow this can be measured using cross-entropy

- note: cross-entropy \geq entropy

- computing per-word cross entropy: $H_n(w_1, \dots, w_n) = -\frac{1}{n} \log_2 \hat{P}_n(w_1, \dots, w_n)$

- perplexity: $2^{(\text{cross-entropy})}$

- interpretation: the average branching factor at each decision point, if our dist. were uniform

- the "goodness" of different values of these measures depends on the corpus. e.g., if corpus is "meow, meow, meow, ..." then we'd hope cross-entropy would be very low (0).

- Smoothing: fixes flaw of MLE, which is that it estimates probs that make training data maximally probable by making everything else (unseen data) minimally probable

- add-one (Laplace) smoothing: just pretend you've seen everything one more time than you did

$$P_{+1}(w_i | w_{i-2}, w_{i-1}) = \frac{C(w_{i-2}, w_{i-1}, w_i) + 1}{C(w_{i-2}, w_{i-1}) + V} \quad \text{where } V = \text{vocab size}$$

- large vocabs make it so that way too much prob mass is stolen from seen events

- add- α (Lidstone) smoothing improves things:

$$P_{+\alpha}(w_i | w_{i-2}, w_{i-1}) = \frac{C(w_{i-2}, w_{i-1}, w_i) + \alpha}{C(w_{i-2}, w_{i-1}) + \alpha V} \quad \text{where } \alpha < 1$$

- to find α : test α on different validation sets to see which gives lowest cross-entropy

n-gram language models

- good-Turing
 - previous methods change the denominator, which can have big effects on frequent events; Good-Turing changes the numerator
 - usually, we have $P_{ML} = \frac{c}{n}$ \rightarrow # times you see n-gram

but Good-Turing uses adjusted counts c^* : $P_{GT} = \frac{c^*}{n}$ where $c^* = (c+1) \frac{N_{c+1}}{N_c}$,

$N_c = \#$ n-grams that appear exactly c times in corpus

$N_0 = \#$ unseen n-grams

• summary:

- add-1 and add- α are simple, but not very good
- Good-Turing more sophisticated and better, but there are even better ones

more smoothing and the noisy channel model

- Good-Turing says these have same prob: "Scottish beer drinkers" and "Scottish beer eaters"
- solution: use info from lower-order N-grams "beer drinkers" and "beer eaters"
- two ways: interpolation and backoff
- interpolation: combine lower order N-gram models since they have different strengths/weaknesses
 - just a weighted average of probs from different models
 - called a mixture model
 - weights λ_i are interpolation parameters or mixture weights
- back-off: trust the highest order language model that contains N-gram
- but diversity of histories matters: think $P(\text{York}|\text{New})$ is high but $P(\text{York}|\text{!New})$ is low
- Kneser-Ney Smoothing takes diversity of histories into account
 - replace raw counts with counts of histories:

$$P_{ML}(w_i) = \frac{c(w_i)}{\sum_w c(w)} \Rightarrow P_{KN}(w_i) = \frac{N_{1+}(\bullet w_i)}{\sum_w N_{1+}(\bullet w)}$$

$$\text{where } N_{1+}(\bullet w_i) = |\{w_{i-1} : c(w_{i-1}, w_i) > 0\}|$$

- best smoothing for n-grams

• quick summary:

- uniform probs: add- α , Good-Turing
- probs from lower-order n-grams: interpolation, backoff
- Prob of appearing in new contexts: Kneser-Ney

- now dealing with word similarity, e.g., knowing that $P(\text{salmon}|\text{caught two})$ tells us something about $P(\text{swordfish}|\text{caught two})$

- can use embeddings

- noisy channel model: $P(Y) \Rightarrow P(X|Y) \Rightarrow P(X)$

Application	Y	X
Speech recog.	spoken words	acoustic signal
machine translation	words in L_1	words in L_2
spelling correction	intended words	typed words

more smoothing and the noisy channel model (cont'd)

- we want $\operatorname{argmax}_y P(y|x)$

$$\operatorname{argmax}_y P(y|x) = \operatorname{argmax}_y \frac{P(x|y) P(y)}{P(x)} = \operatorname{argmax}_y P(x|y) P(y)$$

Spelling correction, edit distance, and EM

- spelling correction:

- we don't know $P(y)$, so need to make constraints:
 - words must differ by one character from x
- then can compute $P(x|y) P(y)$
- assume substitutions, e.g., $o \rightarrow e$, has one prob instead of a prob conditional on context (surrounding letters)

so:

$$P(x|y) = \prod_{i=1}^n P(x_i|y_i)$$

e.g., $P(no|not) = P(n|n) P(o|o) P(-|t)$

data needed:
subs probs

- can estimate subs probs with data (confusion matrix)
- alignments and edit distance
 - want to find optimal character alignment b/t two words (fewest changes: min edit distance MED)

STALL
TALL
TABL
TABLE

$$\text{MED}(\text{stall}, \text{table}) = 3$$

also: STALL -
(as alignment) d l l s l i
- TABLE

- exponential number of possibilities: how to choose?
- use dynamic programming (memoization) algorithms: Viterbi, CKY
- intuition: $D(\text{stall}, \text{table})$ must be min of:

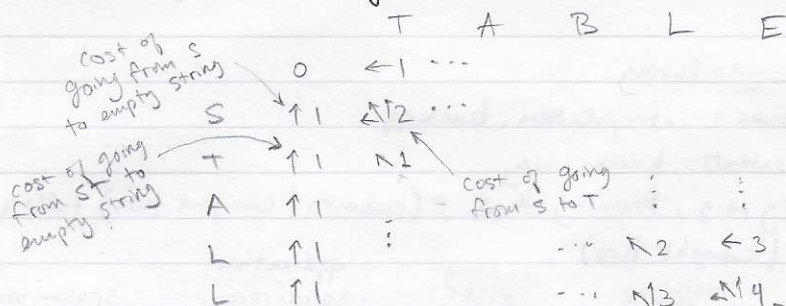
$$D(\text{stall}, \text{tabl}) + \text{cost}(\text{ins})$$

$$D(\text{stal}, \text{table}) + \text{cost}(\text{del})$$

$$D(\text{stal}, \text{tabl}) + \text{cost}(\text{sub})$$

$$\text{assume cost}(\text{ins}) = \text{cost}(\text{del}) = 1, \text{cost}(\text{sub}) = 2$$

- chart stores two things: MED and backpointers



fill like this: start etc.

- chicken-and-egg: want to use costs from noise model, but need costs to estimate noise model

- solution: EM algorithm

- initialize params to arbitrary values (e.g., all costs = 1)
- using these params, compute optimal values (run MED to get alignments)

- using alignments, recompute params
- repeat 2 and 3 until params stop changing

FNLP

final exam notes

spelling correction, edit distance, and EM (cont'd)

→ to something, i.e. doesn't diverge or oscillate

- EM just described for spelling correction is "hard" EM, which converges (like soft EM) but is not nicely defined mathematically (doesn't converge to optimum of likelihood function like soft EM does); but it probably works fine in practice (easier to compute)

- likelihood function:

$$P(\text{data} | \theta) = \prod_{i=1}^n P(x_i | y_i) = \sum_a \prod_{i=1}^n P(x_i | y_i, a)$$

→ if alignments a are latent

- neither hard nor soft EM guaranteed to converge to global optimum; soft may be not even local

text classification

- have y values like spam/not spam, sentiment, topic
- bag of words models: word order doesn't matter
- Naive Bayes

$$\hat{c} = \underset{c \in C}{\operatorname{argmax}} P(c | d) = \underset{c \in C}{\operatorname{argmax}} P(d | c) P(c)$$

model $P(d | c)$ as set of features (words) it contains: $P(d | c) = P(f_1, f_2, \dots, f_n | c)$
 $\approx P(f_1 | c) P(f_2 | c) \dots P(f_n | c)$ (naive Bayes assumption)

So: $\hat{c} = \underset{c \in C}{\operatorname{argmax}} P(c) \prod_{i=1}^n P(f_i | c)$

- can estimate $P(f_i | c)$ with simple smoothing

- alternative feature values and feature sets:

- use only binary values for f_i
- only use subset of vocab for F
- use more complex features (e.g., bigrams, syntactic features)
- for sentiment analysis: whether word is +/-

- quick and easy model

- but independence assumption may be strong / unrealistic

- MaxEnt better if you have enough data

- a.k.a. multinomial logistic regression

$$\hat{c} = \underset{c \in C}{\operatorname{argmax}} P(c | d) \quad (\text{same as NB})$$

... but we model $P(c | d)$ directly

- features are functions of both observations \vec{x} and class c

$$P(c | \vec{x}) = \frac{1}{Z} \exp\left(\sum_i w_i f_i(\vec{x}, c)\right)$$

(logistic)

→ and using softmax

⇒ as if just running c separate regressions...

- training:

$$\hat{w} = \underset{\vec{w}}{\operatorname{argmax}} \sum_j \log P(c^j | x^j)$$

can take some time (gradient ascent): conditional MLE (CMLE)

Part-of-speech tagging and HMMs

- first step toward syntactic analysis
- deciding on correct label depends on word to be labeled and labels of surrounding words
- parts of speech: two main classes: open class ("content" words) and closed class (function words)
- commonly 40-100 different tags
- hard b/c of ambiguity, sparsity
- relationship to previous models:

	sequences	hidden vars
n-gram	✓	✗
NB	✗	✓
HMM	✓	✓

- n-gram model: a model for sequences that also makes a Markov assumption, but has no hidden vars
- Naive Bayes: a model w/ hidden vars (classes) but no sequential dependencies
- HMM: a model for sequences with hidden variables
- HMM:

$$\arg \max_T P(T|S) = \arg \max_T P(S|T) P(T)$$

$$P(T) = \prod_i P(t_i | t_{i-1})$$

$$P(S|T) = \prod_i P(w_i | t_i)$$

$$\text{where } P(S|T) P(T) = \prod_i P(w_i | t_i) P(t_i | t_{i-1}) = P(S, T)$$

- too many tag sequences T to enumerate
 \Rightarrow use Viterbi again (see next lecture)

algorithms for HMMs

- change of notation:

A : transition matrix w/ possible states q^1, q^2, \dots, q^N for N tags

B : output matrix w/ possible outputs o^1, o^2, \dots, o^V for V words

$\lambda = (A, B)$: params of HMM

Q : state sequence q_1, q_2, \dots, q_T } T : time

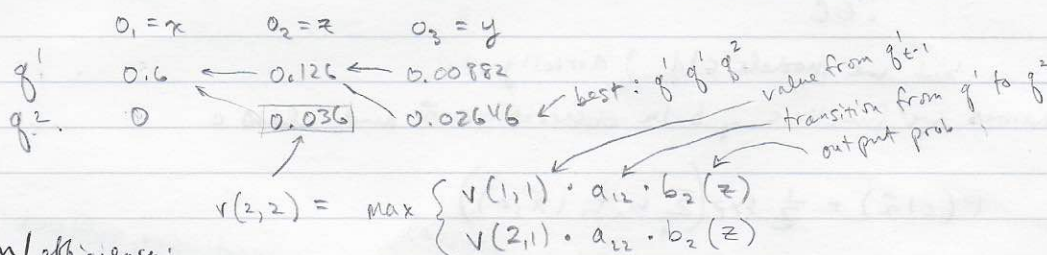
O : output sequence o_1, o_2, \dots, o_T

- using new notation: $w_i \quad t_i \quad t_i \quad t_i \leftarrow \text{old notation}$

$$P(O, Q | \lambda) = \prod_{t=1}^T P(o_t | q_t) P(q_t | q_{t-1})$$

$$= \prod_{t=1}^T b_{q_t}(o_t) a_{q_{t-1}q_t}$$

- example grid:



- implementation/efficiency:

- enumeration is $O(N^T)$, $O(T)$ space
- Viterbi is $O(N^2T)$, $O(NT)$ space (much better on run time)
- use neg log probs

- may also want to compute likelihood: $P(O|\lambda) = \sum_Q P(O, Q|\lambda)$

- use Viterbi (w/ a slightly different update), sum last column for $P(O|\lambda)$

- learning $\lambda = (A, B)$: use EM if unsupervised (no tag labels), using expected counts in E step
- Forward-backward algo: avoid enumerating all possible sequences \Rightarrow an instance of EM

FNLP

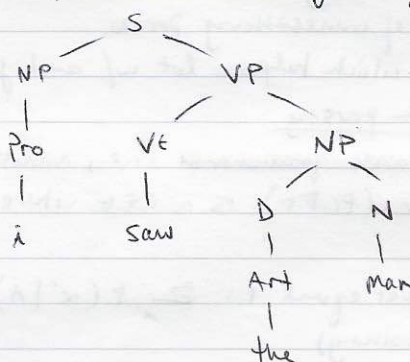
Final exam notes

Methods in annotation and evaluation

- Annotation not easy, even for humans because of ambiguity and gray areas (both syntax and semantics)
- Penn Treebank has 36 POS tags (excluding punctuation) w/ extensive guidelines for annotators
 - even so, there will still be errors
- inter-annotator agreement (IAA)
 - have multiple people annotate, compute raw agreement rate
 - very low rates \Rightarrow examine sources of disagreement and consider rewriting guidelines
 - agreement rate can be considered upper bound (human ceiling) on accuracy of a system evaluated on that dataset
 - can also weight errors by severity
- can make hypotheses about linguistic objects (e.g., author), output of a language system, and human beings
- can use cross-validation to choose hyperparams (e.g., for smoothing)
- measures of performance:
 - precision: $P = \frac{TP}{TP+FP}$ \Rightarrow fixes for cases where you over-predict 1 when there are a lot of 1's
 - recall: $R = \frac{TP}{TP+FN}$ \Rightarrow fixes for cases where you over-predict 0 when there are few 0's
 - F-score: $F_1 = \frac{PR}{P+R}$
- use significance tests to determine whether measurements are different

Syntax and parsing

- theory of syntax should explain which sentences are well-formed (grammatical); well-formed distinct from meaningful
 - two theories: context-free grammar and dependency grammar
- context-free grammar (CFG)
 - two types of grammar symbols: terminals (t): words; and non-terminals (NT): phrase categories like S, NP, VP, PP
 - has rules of the form $NT \rightarrow \beta$, where β is any string of NT's and t's
 - e.g.,



- structural ambiguity: caused by POS ambiguity; e.g., I saw her duck
- attachment ambiguity: e.g., I saw the man with the telescope.
- all parsers have two fundamental properties
 - directionality: sequence in which structures are constructed (top-down, bottom-up, mixed)
 - search strategy: the order in which the search space of possible analyses is explored: depth-first, breadth-first, best-first

example: the dog bit	Step	op	subgoals	input	operations:
(recursive descent strategy) \Rightarrow top-down, depth-first	0			the dog bit	
	1	E	NP VP	"	E = expand
	2	E	DT NN VP	"	M = match
	3	E	the NN VP	"	
	4	M	NN VP	dog bit	Bn = backtrack to step n
	5	E	bit VP	"	
6	B4		NN VP	"	

data: need POS tags and rules

problems: can be many backtracks or infinite loops
(NP \Rightarrow NP PP)

7

Syntax and parsing (cont'd)

- shift-reduce: depth-first, bottom-up

— e.g., —→

- depth-first parsers very efficient for unambiguous structures but massively inefficient when faced with local ambiguity

- but can use probabilistic model to learn which choices to make (next lesson...)

- breadth-first search using dynamic programming

- avoid re-analyzing any substring because its analysis is independent of the rest of the parse
 - memoized in chart, a.k.a., well-formed substring table (WFST).

- CKY algorithm: bottom-up, breadth-first

- takes $O(Gn^3)$, where G is the number of grammar rules

	1	2	3	4
0	Pro, NP ①			S ⑤
1		VE, VP, N	②	VP ④
2			Pro, Pos Pro, ①	NP ③
3				N, Vi
	he ₁	saw ₂	her ₃	duck ₄


① $NP \rightarrow Pro$

② $D \rightarrow Pos Pro$

③ $NP \rightarrow D N$

④ $VP \rightarrow V + NP$

⑤ $S \rightarrow NP VP$



- start with all POS allowed for word

- ignore lower-diagonal: would be backward strings

- to fill cell (i, j) , we use cell from row i and a cell from column j

⇒ So need to fill cells down and left of (i, j) before filling (i, j)

- even though avoids re-computing substructures, much more efficient than depth-first (in worst case)

- but still may be a lot of unnecessary parses.

- next: statistical parsing, which helps a lot w/ ambiguity and efficiency

CKY parsing, treebanks and Statistical parsing

- for probabilistic parsing, use treebank grammars (i.e., annotated sentences)

- a probabilistic context-free grammar (PCFG) is a CFG where each rule $A \rightarrow \alpha$ is assigned a probability $P(\alpha | A)$

- sum over expansions of A must equal 1: $\sum_{\alpha} P(\alpha | A) = 1$

- use MLE for probs (w/ smoothing)

- a generative model, like with HMMs

- prob of parse t is product of all rules of parse: $P(t) = \prod_{A \rightarrow \alpha \in t} P(\alpha | A)$

- we also have a language model since t implicitly includes words \vec{w} : $P(t) = P(t, \vec{w})$

- sentence prob is then obtained by summing over $T(\vec{w})$, the set of valid parses of \vec{w} :

$$P(\vec{w}) = \sum_{t \in T(\vec{w})} P(t, \vec{w}) = \sum_{t \in T(\vec{w})} P(t)$$

i.e. each non-terminal node has a prob; just multiply these together

- straightforward to extend CKY parsing to probabilistic case

- goal: return highest prob parse of sentence (analogous to Viterbi)

- best-first parsing help by not having to do exhaustive parsing

- constituents have scores and are added to agendas, which are ordered by scores

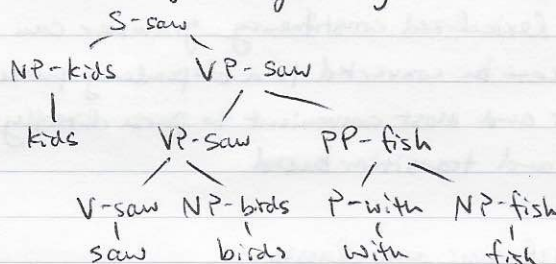
- scores computed as an average over words to normalize tree sizes

- not-so-great consequence: words w/ same POS will give same scores/probs even though we know this isn't true

heads, dependency parsing

ways to fix PCFGs

1. automatically create new categories that include the old category as the parent
 * e.g., an NP in Subject position becomes NP^S
2. lexicalization: create new categories by adding the lexical head of the phrase
 * e.g.,




- but this leads to huge grammar blowup and very sparse data
- evaluating parse accuracy
 - output considered correct if there's a gold constituent that spans the same sentence positions
 - use precision/recall/F₁ score
- dependencies


transforming constituency → dependency parse

- start w/ lexicalized constituency parse, remove phrasal categories, remove duplicated terminals, and collapse chains of duplicates

- so the lexicalized constituency parse above becomes:




- another way of showing this:



- some treebanks prefer content heads, while others functional heads

- e.g., "...were always watching..."

- edge labels:



- projectivity: parse said to be projective if every subtree (node and all its descendants) occupies a contiguous span of the sentence

- i.e., no crossing edges

- nonprojectivity is rare in English, but quite common in other languages

- constituency → dependency parse (again): but how do we find each phrase's head in the first place?

- use head rules: designate one RHS nonterminal as containing the head

direct dependency parsing:

transition-based parsing

- adapts shift-reduce methods: stack and buffer (note: shift-reduce more efficient than CKY)
- idea: train classifier to predict next action (SHIFT, REDUCE, ATTACH-LEFT, ATTACH-RIGHT), and proceed left-to-right through sentence. $O(n)$ time complexity!
- only finds projective trees

graph-based parsing

- global algorithm: from the fully connected directed graph of all possible edges, choose the best ones that form a tree

- e.g., maximum spanning tree algorithm ($O(n^2)$)

heads, dependency parsing (cont'd)

- **conversion-based parser**: first constituency parse, then convert to dependencies
- **Summary**:
 - while constituency parses give hierarchically nested phrases, dependency parses represent syntax with trees whose edges connect words in a sentence
 - head rules govern how a lexicalized constituency grammar can be extracted from a treebank, and how a constituency parse can be converted to a dependency parse
 - for English, often fastest and most convenient to parse directly to dependencies; two main paradigms: graph-based and transition-based

lexical semantics: word senses, relations, and classes

- why it's hard: different senses, synonyms, ^{animal} hyponym/hypernym (subset/superset), similarity and gradation, need for inference
good and great
vacation/holiday
- **WordNet** a hand-built resource w/ 117,000 synsets: sets of synonymous words
bank: river / finance
multiple senses
 - although this doesn't solve issues w/ homonyms, poly senses
 - Synsets organized into a network by several kinds of relations, including:
 - **hyponymy** (is-a): hyponym {ambulance} is a kind of hypernym car
 - **meronymy** (part-whole): meronym {air bag} is a part of holonym car
 - nouns have on average 1.24 senses, verbs 2.17
 - incomplete: doesn't have ^{all} multiword phrases (stress out), neologisms (facepalm), names (Microsoft)
- **word sense disambiguation (WSD)**
 - there are competitions held every 1-3 years to do this
 - data-driven methods do well
 - e.g., Naïve Bayes, decision lists, decision trees
 - as features, can use content words in some window; can also use syntactically related words, syntactic role in sense, topic of the text, POS tag, surrounding POS tags
 - evaluation: extrinsic, intrinsic, baseline (choose most frequent sense - sometimes hard to beat)
 - issues: not clear how fine-grained to be; difficult expensive to annotate w/ fine-grained; classifiers must be trained separately for each word
- **named entity recognition and supersense tagging**: categories like person, location, artefact
for extracting proper names
 - supersenses (e.g., attribute, food, object, quantity, motion, social) are language-neutral

distributional lexical semantics

- **distributional hypothesis**: a word is the company it keeps
- use word-context matrices (like in NLU): each word is a vector of context word frequencies/probs
- two kinds of co-occurrence between two words: first order (near each other), second order (similar neighbors)
- for syntactic similarity, use small context windows; for semantic, use larger (sometimes whole document)
- **collocation**: occurring unusually often in the context of word w
 - can use $PMI(x,y) = \log_2 \frac{F(x,y)}{P(x)P(y)}$
 - b/c uses MLE, it's over-sensitive to chance co-occurrence of infrequent words
 - alternatives; t-test, χ^2 , others...
 - or can improve PMI: positive PMI (PPMI) = $\max(PMI, 0)$
- measuring similarity: cosine: $\cos(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|}$
- sparse \rightarrow dense: SVD/LSA, skip-grams, etc.
 - e.g., brown clusters: binary tree

FNLP

final exam notes

distributional lexical semantics (cont'd)

- class-based language model: $P(w_i | w_{i-1}) = P(c_i | c_{i-1}) P(w_i | c_i)$

Semantic role labeling and argument structure

- can't just use syntax to trivially infer semantic roles
- two computational datasets/approaches that describe sentences in terms of semantic roles:
 - PropBank: simpler, more data
 - FrameNet: richer, less data
- PropBank has structure and its roles (Arg0, Arg1, ...)
 - abstracts away from syntax to predicate-argument structures
 - lexicon and annotations on full WSJ PTB corpus and other data
 - limitation: roles (Arg0, ...) are predicate-specific
 - originally verbs only, but now has other POS
- FrameNet → like PropBank
 - no longer (primarily) verb-based, but frame-based
 - FrameNet parsing: SEMAFOR
 - pipeline similar (but not identical) to traditional SRL

machine translation [non-examinable]